



Automated vehicle's behavior decision making using deep reinforcement learning and high-fidelity simulation environment[☆]

Yingjun Ye, Xiaohui Zhang, Jian Sun*

Department of Traffic Engineering & Key Laboratory of Road and Traffic Engineering, Ministry of Education, Tongji University, Shanghai, China



ARTICLE INFO

Keywords:

Automated vehicle
Decision making
Deep reinforcement learning
Reward function

ABSTRACT

Automated vehicles (AVs) are deemed to be the key element for the intelligent transportation system in the future. Many studies have been made to improve AVs' ability of environment recognition and vehicle control, while the attention paid to decision making is not enough and the existing decision algorithms are very preliminary. Therefore, a framework of the decision-making training and learning is put forward in this paper. It consists of two parts: the deep reinforcement learning (DRL) training program and the high-fidelity virtual simulation environment. Then the basic microscopic behavior, car-following (CF), is trained within this framework. In addition, theoretical analysis and experiments were conducted to evaluate the proposed reward functions for accelerating training using DRL. The results show that on the premise of driving comfort, the efficiency of the trained AV increases 7.9% and 3.8% respectively compared to the classical adaptive cruise control models, intelligent driver model and constant-time headway policy. Moreover, on a more complex three-lane section, we trained an integrated model combining both CF and lane-changing behavior, with the average speed further growing 2.4%. It indicates that our framework is effective for AV's decision-making learning.

1. Introduction

The automated vehicles have captured the public attention in recent years, especially after Google announced its automated driving program in 2010, for its advantages of alleviating the traffic congestion, liberating drivers' attention and conserving energy. The tasks involved in achieving autonomous driving can be divided into three modules: environment recognition, decision making and vehicle control. Among them, the vehicle control has no obvious differences between AV and human driving vehicle. Meanwhile, the other two modules both turn to deep learning methods for better solutions in recent development. The AI-based environmental recognition methods including convolutional neural network (CNN) (LeCun et al., 1998) and deep belief network (Hinton et al., 2006) are used to map from an image to perception information, such as detection of traffic participants (Szarvas et al., 2005; Wang et al., 2014), lane markings (McCall and Trivedi, 2004; Fritsch et al., 2013), traffic signs (De La Escalera et al., 1997; Sermanet and LeCun, 2011), etc. Studies on end-to-end learning integrating three modules are also concerned in the autonomous driving field, that is to map raw pixels from a camera directly to steering commands (Bojarski et al., 2016). However, studies specified to decision-making task are very limited for various reasons. First, a large quantity of field data is the basis of training with deep learning algorithms, while data collection is a high-cost and time-consuming job. Many datasets, like KITTI (Geiger et al., 2013) and ImageNet (Deng et al., 2009), are used for training environment recognition, while no public data available for decision-making behaviors.

* This article belongs to the Virtual Special Issue on "CAV Control".

Corresponding author.

E-mail address: sunjian@tongji.edu.cn (J. Sun).

Secondly, it is hard to label the decision-making behavior, given that actions are not only stimulated by the surrounding environment but also by the drivers' psychological factors (Van Winsum, 1999), while the latter cannot be easily observed and quantified. What's more, traditional imitation learning algorithms limit the operation skill to human understanding (Kober and Peters, 2010). That means, supervised learning from human driving data, can only imitate the driving behavior of sampled drivers, but not necessarily the optimal driving skills. And the generalization ability to a new environment is also doubtful.

On the other hand, deep reinforcement learning (DRL) algorithm is proved to perform excellently on decision-making tasks. It solves the problems of relying on historical data and imitation to the human player, the resulting agent might act much better than a human does. Google's DeepMind team first successfully applied DRL to learn policies for agents on Atari 2600 games (Mnih et al., 2013, 2015, 2016; Lillicrap et al., 2015), with the results achieved even better performance than an expert human player. The final revolutionary achievement was the AlphaGo, an artificial intelligence computer program plays the Go, a board game that might be humankind's most complicated one. It indicates that DRL could achieve better performance than human beings in highly complex tasks.

It is known that three elements are critical to DRL: (1) Training environment. It should contain a variety of scenes to improve the generalization. And the rivals or interactive objects should be experts or good players at least, so that agent can learn from good examples. (2) Reinforcement learning algorithm. Each case corresponds to its own proper algorithm. For example, Q-learning is a simple but effective algorithm which can be only adopted in the discrete-state problem. (3) Reward function. The training objective is incorporated into reward function to evaluate the reward or punishment of every action. However, the existing studies on autonomous driving have some insufficiencies in all these three aspects. The racing games (e.g. TORCS, GTA-V) were chosen as the training environment whose traffic flow scenes are not realistic. The input states and output actions are discretized, which seriously affects the accuracy of prediction. Furthermore, though deep RL is confirmed to be unstable and difficult to learn effectively in practice (Henderson et al., 2017; Ilyas et al., 2018; Islam et al., 2017; Mania et al., 2018), there are no uniform methods of reward function setting and parameter tuning.

Therefore, to improve the practicability of the deep RL on autonomous driving, based on a more realistic driving environment, we study the modeling of decision-making behavior with a modified DRL model in this paper. More specifically, there are mainly three contributions of this paper: (1) An integrated framework based on VISSIM high-fidelity simulation and DRL is proposed for decision-making training. (2) Three attempts of reward functions are investigated to explore their influence mechanism. What's more, the methods of verifying the effectiveness and two ways of accelerating convergence are discussed. (3) The proposed methods are first tested to train the CF model on a one-lane freeway section, and then extended to a three-lane section to learn the comprehensive model integrates CF and LC behavior. Results indicate that it is effective to learn multi-task driving simultaneously with our platform.

The rest of paper is structured as follows. Section 2 reviews the recent works about the decision making of AV. Section 3 introduces the framework of our training platform and each component in detail. A case study of the CF behavior is presented in Section 4, and discussion is made on reward setting. An extended model integrates LC behavior is also presented. The last section presents the conclusion and future work.

2. Works related to AV's decision making

Though future scenarios will involve mixed flow with AVs and human-driven vehicles, or even 100% AVs, however, the first step of autonomous driving is individual driving under the multi-manually driving traffic flow environments. For a step-by-step development to fully automation, both researchers and manufacturers mostly concentrate on solution of single AV at current stage. Therefore, the work on decision making is mainly about single AV's in this section.

The most influential competition of AVs is possibly the Grand Challenge launched by the (DARPA). The first two competitions (2004, 2005) both conducted on one-way desert trails, thus more attention was paid on environment recognition and vehicle control to deal with different road conditions. Until these two modules were well-learned, the more complicated decisions making were involved in the third event, the Urban Challenge (2007). It called for AVs to drive through a much simplified urban environment, interacting with other vehicles and obeying the traffic rules. However, for getting higher rank, the competitors mostly used rule-based approaches (Buehler et al., 2009) to ensure absolute security, which was kind of conservative. Though great advances had been made compared to the first competition, the state of technology was still far from practical use.

It is generally known that simulation could offer simpler and less expensive alternatives, and allow direct comparison between agents under identical conditions. Therefore, after the DARPA Grand Challenge, to improve the decision-making behavior, international game AI competitions (WCCI-2008, CIG-2008, CEC-2009 and GECCO-2009) were held and drawn the public attention from field test to virtual simulator training. The Open Racing Car Simulator (TORCS) (Torcs, 2017) was one of the competition programs. Hand-crafted rule-based approaches were adopted in early research but resulted in limited performance (Buckland; Chen et al., 2015). To make it more efficient, researchers turned to artificial intelligence to find the solution. Driving behaviors were controlled with general machine learning approaches (Togelius et al., 2006), like neural network and k-nearest neighbor, but still did not prove to give good results. Accordingly, The Neuro Evolution of Augmenting Topologies, tuning both the weighting parameters and structures of networks to find a balance between the fitness of evolved solutions and their diversity, had been applied to improve the performance of original neural network (Togelius et al., 2007; Cardamone et al., 2009a, 2009b). Other improvements of the machine learning approaches included enhancing the generalization capabilities of the evolved model (Kim et al., 2012), optimizing to multi-objectives (Van Hoorn et al., 2009), etc. These machine learning algorithms are able to learn driving to some extent, however, successful applications need large amounts of hand-labeled training data. Another issue is that those supervised learnings on human-driving data are limited to imitation to human drivers, which is able to drive but performed not necessarily excellent.

Until Google's DeepMind team introduced their DRL into behavior prediction, a totally new idea came into researchers' view. Since then, much similar work (Loiacono et al., 2010; Koutník et al., 2014; Lau, 2016; Sallab et al., 2017) had introduced DRL into TORCS driving. Meanwhile, some other games with more realistic graphics and other cars and pedestrians presented were also used to be the training platform, such as World Rally Championship 6 (WRC6) (Richter et al., 2016) and Grand Theft Auto V (GTA V) (Perot et al., 2017), but the results were still not satisfactory. They all focused on the single controlled vehicle and learned to keep in lanes merely, with few interactions with other traffic participants. The much more complex interaction with nearby vehicles on adjacent lane was rarely considered. Furthermore, the traffic scenes in those games only reproduce high-fidelity objects, while the characteristics of traffic flow were not really simulated.

3. Methodology

Generally, there are two types of training methods for decision making modeling (e.g. car-following modeling, lane-changing modeling). The first yet commonly-adopted one is to use empirical microscopic traffic flow data (e.g. real traffic trajectory data or simulator data) to calibrate an analytical or data-driven model. However, for modeling behaviors of AV, the real-world data are not enough and accurate for complex traffic environment. More specifically, massive work on data collection and labeling are the biggest limitation of the use of real-world data (Farabet et al., 2013; Chen and Lin, 2014; Najafabadi et al., 2015). The model trained with data collected in one specific traffic condition is not possible for being perfectly applied on other conditions. Hence, we use another online training method (deep reinforcement learning) based on VISSIM in this paper. The principal is to train the decision-making model from the interaction between "gradually-growing" AV and surrounding VISSIM vehicles. The online VISSIM data does not provide any learning instructions but only perform as the "realistic" traffic environment while countless scenarios can be easily generated in VISSIM. AV can learn driving skills to deal with all kind of traffic scenarios based on the platform. In addition, as the first method is to imitate the existing driving behavior and the reinforcement learning learns to optimize higher rewards, even better driving skills than human drivers are expected from the online training method.

3.1. Framework of the training platform

The feasibility of training autonomous driving by DRL was verified in previous studies (Van Hoorn et al., 2009; Loiacono et al., 2010; Koutník et al., 2014; Lau, 2016; Richter et al., 2016; Perot et al., 2017; Sallab et al., 2017). However, the results were far from practical use for its simplifications from the driving environment to training target. Therefore, the training environment of previous studies, racing games, are replaced by traffic simulator VISSIM to provide more realistic interaction between traffic participants. And the DDPG algorithm is adopted which is suitable for the continuous decision spaces of both output states and input actions in autonomous driving.

As shown in Fig. 1, the training platform of autonomous driving in this paper consists of two parts: the virtual simulation environment and the DRL training program. In every training step, the recognition information in VISSIM is passed to DRL program through the driving simulator interface. Then, the DRL program instructs the AV's action in VISSIM according to recognition information and nearby vehicles in VISSIM will also take response to the movement of AV. At the end of the step, the DRL program evaluates the value of the last action by the reward function and updates the network parameters. The updates are done recursively with the steps going on until a feasible solution is found. Data is passed via TCP/IP communication protocol between the VISSIM and DRL program.

3.2. Training environment–VISSIM

VISSIM is a widely used commercial traffic microscopic simulation software developed by the PTV (The developer of VISSIM). It provides various realistic traffic flow scenes including the road infrastructures and behavior of all road users, while the racing games

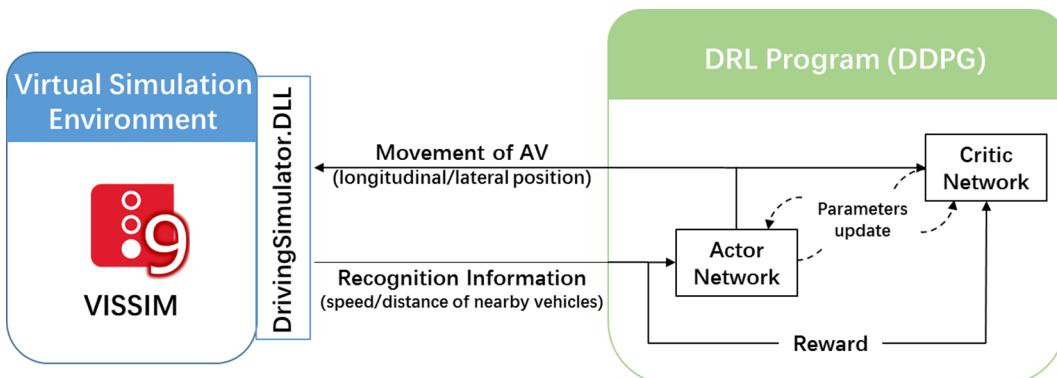


Fig. 1. The framework of the training platform.



(a) Screenshot of GTA-V



(b) Screenshot of VISSIM

Fig. 2. Screen comparison between racing game and traffic simulator.

focus on the 3D visual presentation (see Fig. 2). With rich application programming interfaces, many studies have been made on the virtual testing of AVs in VISSIM.

3.3. Training method—DRL

Reinforcement learning (RL) is defined as a sequential decision-making problem of an agent that has to learn how to perform a task through trial and error interactions with an unknown environment which provides feedback in terms of numerical reward (Van Hoorn et al., 2009). To solve the real problem, theoretical studies have continued to develop since the conception of RL is put forward. Therefore, the following parts firstly introduce the key approaches in the development of RL and their applications. Then the specific theoretical model in this paper is provided in detail.

3.3.1. The development of RL

Markov Process is the theoretical basis of the RL, which can be expressed as a tuple of $\{S, A, P(a|s), \gamma, R\}$, where: S-set of states, A-set of actions, $P(a|s)$ -state transition probability, namely the probability of transiting from state s by taking action a , γ -discount factor, R-reward.

We use the mapping function $\pi(a|s) = P(a|s)$ to denote the policy that agent takes the action a when being in a state s . Also we define a value function to define the expected total discounted rewards got by adopting this policy as Eq. (1), the objective is to find the optimal policy that maximizes the total reward.

$$Q_\pi(s, a) = E_\pi \left[\sum_{k=0}^{\infty} \gamma^k r_k | s, a \right] \quad (1)$$

Q-learning (Watkins, 1989) is perhaps the most classic and commonly used algorithm to solve the MDP problem. It updates the action-value stored in Q-table based on temporal difference incremental learning, which can be expressed as:

$$Q(s_t, a_t) = Q(s_t, a_t) + \alpha [r(s_t, a_t) + \gamma \max_{a_{t+1}} Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)] \quad (2)$$

where, α is the learning rate. It can be only used to solve the MDP problems with actions and states space both discrete and finite. With an initial Q-table and predefined policy, in each step, the agent selects an action a_t based on the current state s_t . According to the reward $r(s_t, a_t)$ feeds back and observation of the next state s_{t+1} , the Q-table is updated by Eq. (2). The steps repeat until Q function converge to an acceptable level.

Loiacono et al. (2010) tried to learn to overtake through Q-learning. They designed two separated scenes: overtaking on a straight sketch and on a tight bend. But to reduce the action space, the agent only learns to steer wheels in the former scene while acceleration and brake are learned in the bend scene, thus the acceleration and steer are not integrated. Moreover, the discrete state and action space cannot precisely describe the interaction between vehicles.

Deep Q Network (DQN). The aforementioned Q-learning cannot deal with the problem with continuous states and does not have generalization ability. For this reason, the Google DeepMind team put forward a DRL structure combines CNN and Q-learning (Mnih et al., 2013, 2015). This method is applied on seven Atari 2600 games, with better performance on six of the games comparing to other previous approaches, and three of them surpass even a human expert.

In short, deep Q network replaces Q-table in Q-learning by the deep neural network (DNN), which compensate the limitation of finite states. The loss function of the DNN is defined as:

$$\text{Loss} = [r_s^a + \gamma \max_a Q^\pi(s', a'; \omega^-) - Q^\pi(s', a'; \omega)]^2 \quad (3)$$

where, ω^- and ω are parameters of target Q network and Q network. Target Q network is a lagging network, the final networks converge when the loss reaches a certain small value, that means the parameters of these two model is quiet close.

3.3.2. Methodology of DDPG

Deep Q learning achieves the input of the continuous state, but the output actions are still discrete, which does not meet the realistic demand. To solve this problem, the DeepMind team further proposed the **DDPG** (Lillicrap et al., 2015). It combines the DQN, the Deterministic Policy Gradient (DPG) and Actor-Critic algorithm.

As Fig. 1 shows, the most revolutionary change is the introduction of the Actor-Critic algorithm. The Actor-Critic algorithm has two neural networks: (1) The critic network $Q(s, a; \omega)$ takes the action from actor network as the input, and outputs corresponding action value. It updates based on the loss of Q network, which is same as the Eq. (3). (2) The actor network uses the DPG to output action policy $P(a|s)$. To maximize the value of the current action, the parameters of actor network σ update by gradient ascent, which uses the chain rule as follows:

$$\frac{\partial Q}{\partial \sigma} = \frac{\partial Q(s, a; \omega)}{\partial a} \times \frac{\partial a(s, \sigma)}{\partial \sigma} \quad (4)$$

With two networks performing their tasks independently meanwhile correlative, not only the efficiency but also the accuracy is improved.

Besides, some other tricks such as experience replay are also incorporated to eliminate the data relevance. The basic idea is that by storing an agent's experiences, and then randomly drawing batches of them to train the network, we can more robustly learn to perform well in the task. By keeping the experience we draw random, we prevent the network from only learning about what it is immediately doing in the environment, and allow it to learn from a more varied array of past experiences. The Experience Replay buffer stores a fixed number of recent memories, and as new ones come in, old ones are removed. When the time comes to train, we simply draw a uniform batch of random memories from the buffer and train our network with them.

Based on DDPG, DeepMind team trained the virtual vehicle in TORCS and achieved excellent driving along the track (Lillicrap et al., 2015). Motivated by the successful demonstrations by DeepMind, some researchers did similar work and optimize the algorithm for handling partial observation and reducing computational complexity (Sallab et al., 2017), and also preliminary discussions are made about different reward functions (Lau, 2016). But these studies merely taught vehicle to drive along the track, the interaction with other traffic participants was not in consideration. What's more important, the reward function, which is actually critical to DRL's training and converge, are not provided with detailed setting methods, what we will discuss in next section.

4. Case study

To testify the feasibility of our training platform and explore the tips about model training, we simplify the scene to single AV's training. And the CF behavior is selected as the first training task. Though CF behavior is deemed to be the most basic decision in driving, it is actually not easy to learn from scratch. With no driving experience and strict restrictions at all, the AV is expected to learn to follow the preceding vehicle with proper distance and steady acceleration. After successfully applied on CF behavior, our training platform is extended to more general driving tasks incorporating LC and overtaking behavior.

4.1. Evaluation criterion

In supervised learning, one can easily track the performance of a model by evaluating it on the validation sets. In reinforcement learning, however, accurately evaluating the training effects can be challenging. Since the training objective is to maximize the total reward the agent collects averaged over a number of games. Thus the total reward or average Q value in one episode, and the loss of the critic network are commonly used in previous studies (Mnih et al., 2013; Lillicrap et al., 2015; Mnih et al., 2015) to evaluate the stability. Their reflections on the training convergence are essentially consistent. Nevertheless, these indexes fluctuate between episodes, that leads to a mathematical imprecise. Researchers only judge the convergence visually by line charts, though it is not objective enough. Therefore, we derive the approximated upper bound of average action value (refer to Appendix A for the concrete process) to help to quantitatively evaluate the stability:

$$\lim_{t \rightarrow \infty} Q_t \approx \frac{r}{1 - \gamma} \quad (5)$$

As for effectiveness, the score is the most intuitive method in previous game training. It is even compared to the score played by the human player to evaluate the intelligence of agent. The environment platform in this paper is a traffic simulator, with no scoring system. Meanwhile, there is no universal standard for autonomous driving evaluation. So, firstly, we evaluate the effectiveness of the trained driving strategy by visual observation. Then, for the model seems "effective", the constant-time headway (CTH, Swaroop and Rajagopal, 2001; Li and Shrivastava, 2002) and intelligent driver model (IDM, Kesting et al., 2007, 2008), two generally accepted adaptive cruise control (ACC) models are used to be the evaluation standard.

4.2. Variable definition

Variables used in the experiments are summarized in Table 1.

Table 1
Symbols and parameters.

Variables and parameters	
d	Net distance from preceding vehicle (m)
v	Speed of the AV (km/h)
Δv	Speed difference to preceding vehicle (km/h), $\Delta v = v - v_{pre}$, normalized by 40 km/h
Δa	Jerk (m/s ³)
h_d	Desired time headway (s), $h_d = 2$
State flags	
C	Collision, $d < 0$ m
L	Low-speed following though preceding vehicle is far away, $d > 15$ m & $v < 1$ m/s & $\Delta v < 0$ m/s
U	Uncomfortable jerk, $\Delta a > 5.6$ m/s ³ , it is the maximum jerk passenger cannot notice
R	Reverse, $v < 0$ m/s
O	Otherwise, vehicle moves in normal state

4.3. Network structure

We have demonstrated in DDPG methodology that it contains two networks: the critic model estimates Q-values while the actor network choose the optimal actions based on these estimations. Specifically, the network structures are presented in Fig. 3. The actor network has 4 layers, those are 1 input layer, 1 output layers and 2 hidden layers with 300/600 units in between. Both hidden layers used ReLU activation function, for its advantage of accelerating the convergence (Maas et al., 2013). The output layer is comprised of tanh activated neurons, which produces continuous bounded action. In the critic model, the action is not made visible until the second hidden layer with other 300 hidden units. The first and third hidden layers are ReLu activated, while the second merging layer computes a pointwise sum of a linear activation computed over the first hidden layer and the action input. The network hyperparameters are shown in Table 2.

4.4. Training tips about model training

It is known that some implementation details are the key factors of the deep learning (Wei, 2015), such as the initialization and fine-tuning methods. However, regarding DRL, the major difference is the reward function, which has not been discussed in previous studies. Another problem of DRL is the consuming time, days of training is common especially for a complex issue. Thus, some training tips about the reward setting and accelerated methods are necessary, which are discussed in this section.

The simulation environment in exploring cases are CF scenes with a 2 km-long single straight lane, with no signal control and other traffic facilities. To provide appropriate and stochastic traffic flow situations, the flow input is set randomly from 600 to 800 veh/h/ln. With the moderate input flow, we can provide more realistic car-following and lane-changing (LC) situations, while lower

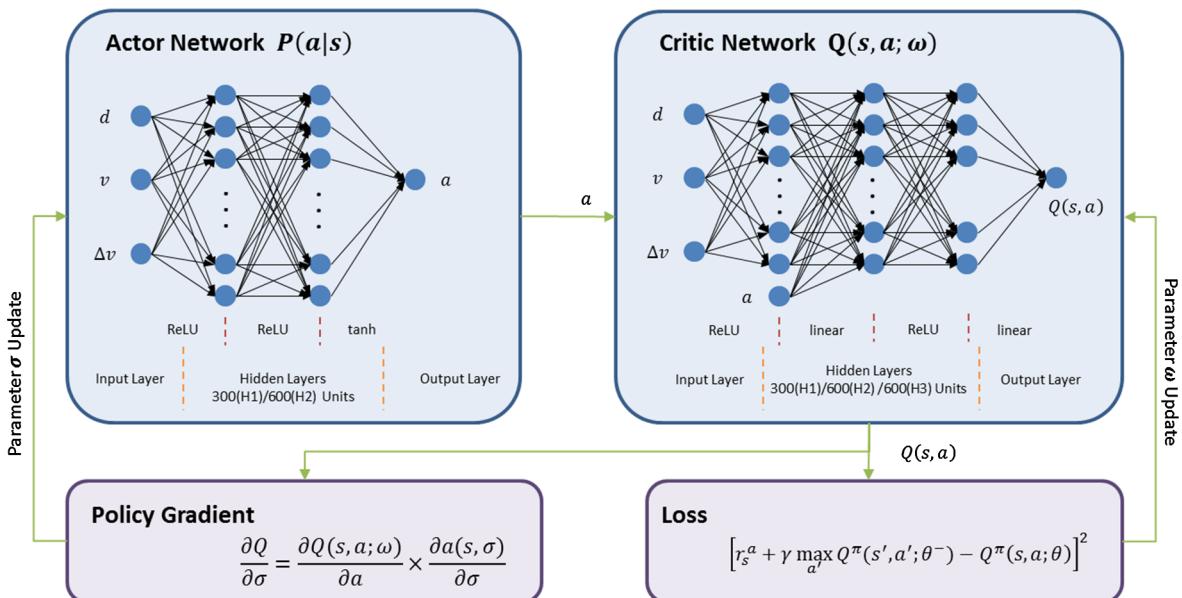


Fig. 3. Actor-critic neural network structures for DDPG.

Table 2
Hyperparameters and its values for DDPG.

Hyperparameters	Value	Descriptions
Discount factor	0.99	Q-leaning discount factor gamma
Update rate	0.001	Update rate of target network tau
LRA	0.0001	Learning rate of actor network
LRC	0.001	Learning rate of critic network
BATCE_SIZE	50	Number of cases used by stochastic gradient descent update
BUFFER_SIZE	100,000	Number of total cases in replay buffer

input flow may easily lead to free-flow driving and heavier flow may easily produce congestion. What's more, the random seed is changed after each episode training. The simulation resolution is 0.1 s. For each episode, the AV comes into the road after 30 s warm-up time. Three situations lead to a new simulation episode, those are the collision, unreasonable stop (far away from the front vehicle) or reaching the end of the road. Since we only study the decision-making module in our study, it is assumed that the environment information has been recognized by sensors, with no noise error. Recognition input is the current speed and acceleration of the AV, the spacing and speed difference to the preceding vehicle, movement information output is the longitudinal acceleration.

4.4.1. Discussion on reward function

The reward function represents the expected driving behavior. For example, it is known that for AV, the desired objectives are safety, high efficiency and driving comfort. Therefore, when the AVs are in unsafe state or drive unstably, the reward would be negative punishment. On the contrary, higher speed will receive a greater reward.

Even following the objectives above, the reward can still be unfeasible or inefficient if the specific function is improper. Therefore, three reward attempts are tested to find the most effective reward form. At first, two rewards with addition form and multiplication form are investigated respectively. Though the rewards seem rational, the results are not satisfactory. We find that it is caused by unmatched magnitude between reward and punishment. Thus, the normalized form and 100-time magnified form to normalized one are raised. But the latter cannot converge to a good result, which is finally solved by the regularization. The detailed discussion of each reward form is listed below.

(a) Addition form and multiplication form

Our initial set of experiments use reward functions respectively with additional form r_{a1} and multiplication form r_{a2} , which are listed as below:

$$r_{a1} = \begin{cases} -100 & C \text{ or } R \\ -80 & U \\ -50 & L \\ h_d \times v - d & O \end{cases} \quad (6)$$

$$r_{a2} = \begin{cases} -100 & C \text{ or } R \\ -80 & U \\ -50 & L \\ \frac{v}{(d - h_d \times v) \times \Delta v} & O \end{cases} \quad (7)$$

The punishments are totally same. They forbid collision (C) and reverse (R), the uncomfortable (U) or inefficient driving performances (L) are also panelized. Their rewards are different but similar to expecting higher speed and proper smaller distance. The reward in Eq. (7) is motivated by the constant time headway policies in ACC system, where vehicles are supposed to keep the speed close to the preceding vehicle and the distance proportional to speed. Though both of the reward functions seems to be rational, results are quite different.

According to the final performance, vehicles trained with the addition form r_{a1} always accelerate to the desire speed until collide into the preceding vehicle. Vehicles with r_{a2} are able to keep a proper distance, however, the frequent sharp speed-up and brake make the jerk to a uncomfortable state. Corresponding to the above results, as shown in Fig. 4, the total rewards of r_{a1} are always negative. Though rewards of r_{a2} have reached a relative high positive value, it shows no sign of convergence.

Even tried with different learning rates, no great progress has been presented. By theoretical analysis, we find the local optimal solution could be due to the magnitude gap between the reward value and punishment, which leads to a large step-size.

(b) Normalized form and 100-time magnified normalized form

In the third reward function, we set the magnitude of the reward and punishment to the same level. Since it is difficult to define the boundary of the multiplication form, we normalize the reward of the addition form r_{a1} , the other variant is magnifying the normalized form by 100:

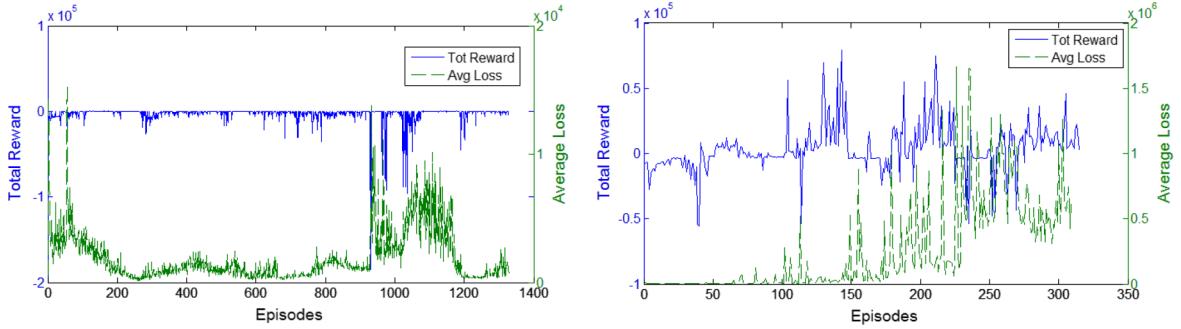


Fig. 4. The total reward and average loss per episode with addition form on the left and multiplication form on the right.

$$r_{a3} = \begin{cases} -1 & C \text{ or } R \\ -0.8 & U \\ -0.5 & L \\ v' - d' & O \end{cases} \quad (8)$$

$$r_{a4} = \begin{cases} -100 & C \text{ or } R \\ -80 & U \\ -50 & L \\ 100 * (v' - d') & O \end{cases} \quad (9)$$

where, d' and v' are normalized variables of d and v , respectively normalized by 100 m and 80 km/h.

Training with r_{a3} , the learning procedure converges after about 800 episodes (see in Fig. 5). What's more, the resulting automated vehicles are able to drive with safety and comfort. In other words, vehicles could keep a proper distance from preceding vehicle, meanwhile avoid frequent or sharp acceleration and brake. Besides, assuming that the expected headway is 2s, according to the upper bound of action value evaluated by Eq. (5):

$$\max(r_{a3}) \approx \frac{80}{80} - \frac{2 * 80/3.6}{100} = 0.56 \quad (10)$$

$$Q_{a3}^{up} = \lim_{t \rightarrow \infty} Q_t \approx \frac{\max(r_{a3})}{1 - \gamma} = \frac{0.56}{0.01} = 56 \quad (11)$$

The maximum of average Q value in our training procedure is 42.0 which is close to the evaluated upper bound. It verifies that normalized form is an effective reward function.

However, Fig. 6 shows that the 100-time magnified form r_{a4} almost learn nothing, the action value even converge to a negative value.

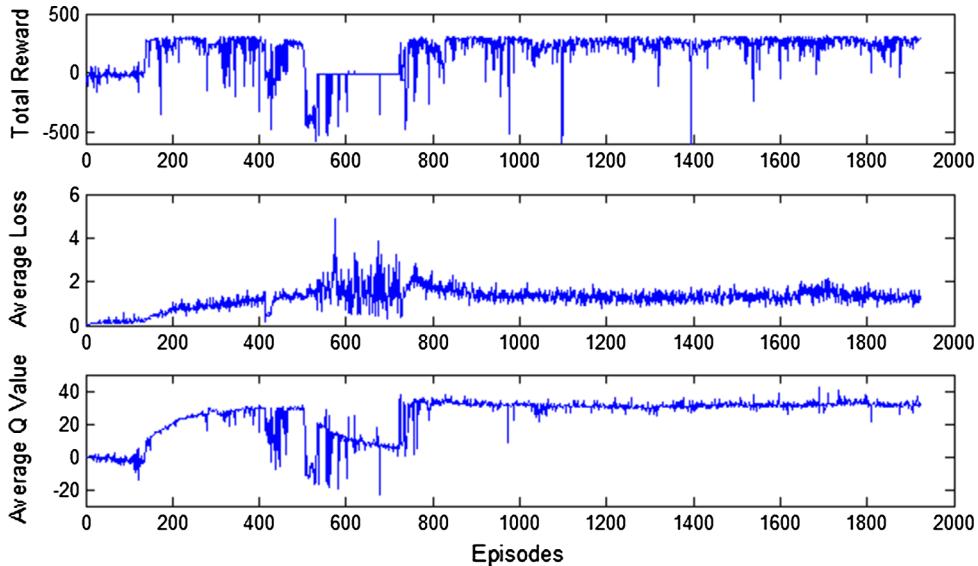


Fig. 5. The evaluation of the normalized reward.

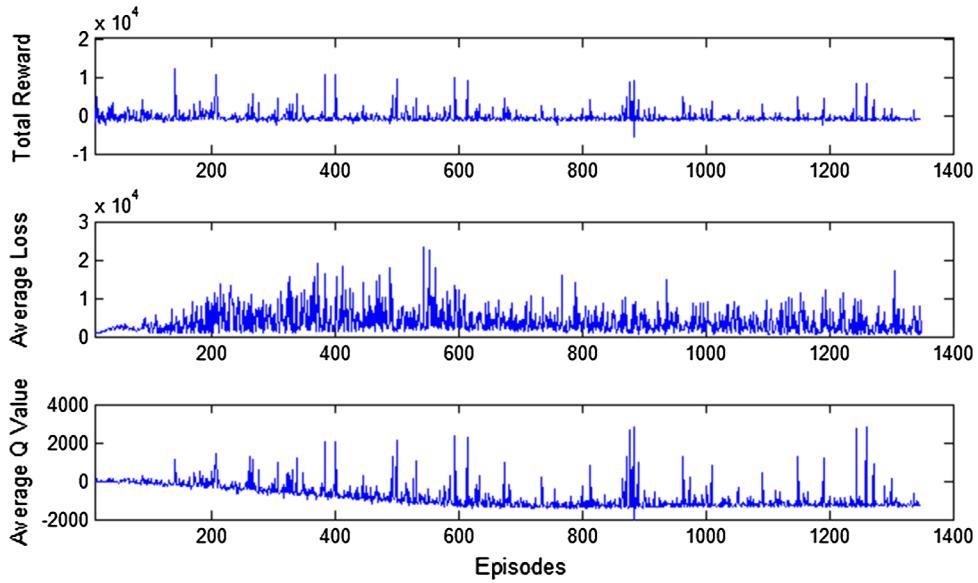


Fig. 6. The evaluation of the 100-time magnified normalized form.

(c) Regularized bounded form

Further analysis of the training mechanism demonstrates that 100-time magnified normalized form almost learns nothing is caused by the “dying ReLU”, the common problem of activation function of critic network. Leaky ReLU is one common method to alleviate this issue, but the consistency of the benefit across tasks is presently unclear (Teaching, 2018). We have tried leaky ReLUs with different slopes, the results are not satisfactory. Since the “dying ReLU” arises with the too high learning rate, naturally, another solution is by tuning the learning rate. According to the parameter ω update algorithm of critic network Q:

$$\omega \leftarrow \omega - \alpha \times \frac{\partial \text{Loss}}{\partial \omega} \quad (12)$$

$$\frac{\partial \text{Loss}}{\partial \omega} = -2 \times \sqrt{\text{Loss}} \times \frac{\partial Q^\pi(s, a; \omega)}{\partial \omega} \quad (13)$$

where α is learning rate. It is clear that the step-size of update parameter is proportional to the square root of loss. Associated with the loss function in Eq. (3), loss is composed of single-step reward and Q value which is the discounted cumulative reward. Hence, the magnification of reward also leads to magnification of loss. Hence, to keep the update step-size to a proper level, the learning rate should be limited to a relatively low level.

On the other hand, the magnification also affects the action network. The activation function of action network, *tanh* function might also fall into local solution when the step-size is too big. Regularize the action network is effective for this problem, that is introducing parameters of action network into loss.

To validate our theory, we still adopt the reward r_{a4} , but minify the learning rate to 1% of original value. At the same time, we regularize the action network. The result shown in Fig. 7 is similar to the normalized form. The model converges after 700 episodes’ training.

Therefore, to make the model effective and quickly converged, there are mainly three key points on setting the reward function: 1. keep the reward and punishment bounded; 2. adjust the learning rate; 3. regularize the action network.

4.4.2. Ways to accelerate training

Generally, due to the high computational complexity, it spends considerable time on training the model. What’s more important, the optimization of network parameters takes dozens of times’ training. Apart from the hardware acceleration, we introduce two accelerated tips by primary trained parameter and enlarged time step.

(a) Based on primary trained parameter

The previous training all “learn from zero”, that is automated vehicle has no driving experience which leads to frequent collision or long following headway. It’s easy to understand that it takes a rather long time from beginner to skilled “driver”. So, we consider whether it is possible to “learn from the experienced” and then get further improvement. For this purpose, firstly we train the action network with the microscopic trajectories collected from VISSIM whose environment is totally same as the DRL training. As shown in Fig. 8, it converges after about 400 episodes, while it takes almost 800 episodes to learn from zero. Therefore, the DRL based on

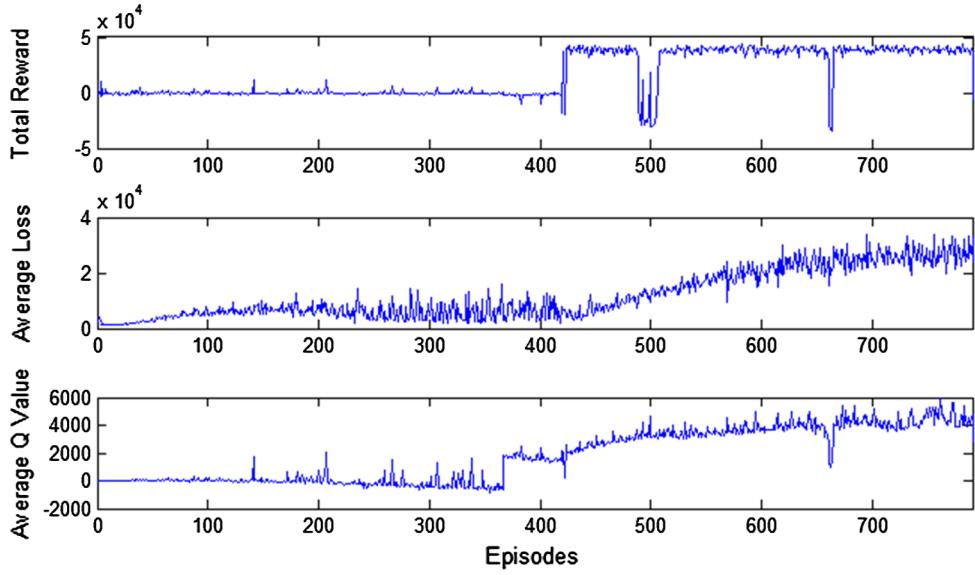


Fig. 7. The evaluation of the regularized bounded form.

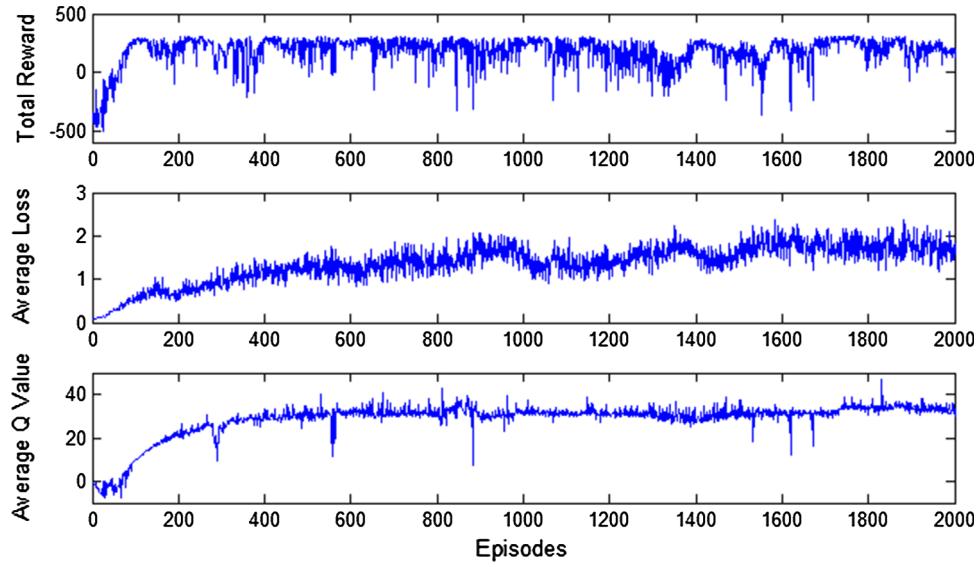


Fig. 8. The evaluation of the normalized form with pre-trained parameter.

primary trained parameter do accelerate training.

(b) Properly enlarge time step

All the time steps in previous cases are set as 0.1 s, that is to say, all the vehicles in the network update their status for every 0.1 s. However, human drivers do not adjust their actions every moment, but continue for some time after every action instead. So, by properly enlarging time step, we can shorten the training time in the premise of assuring effectiveness. Theoretically, if we enlarge the time step to 1 s, it leads to 10-fold increases in efficiency. Results in the Fig. 9 show this method is feasible, the 1 s model converges at about 800 episodes which are similar to the 0.1 s model.

4.5. Evaluation of the DRL trained CF model

On the previous discussion about reward function, we have tentatively verified the effectiveness of the DRL model. To quantify the evaluation of this model, we compare the DRL to IDM and CTH, the widely used CF models for ACC, on the performance of efficiency and comfort. The DRL model is trained with the normalized reward r_{a3} .

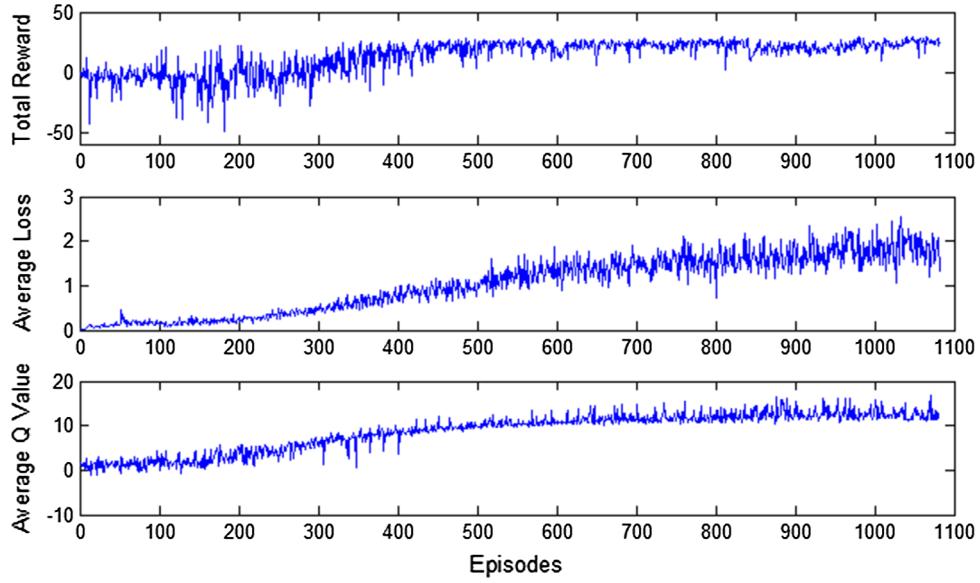


Fig. 9. The evaluation of the normalized form with large time step.

Regarding efficiency, we extract the 800 episodes of DRL and ACC models respectively and assess their average speed for each episode. For DRL model, we take the last 800 converged episodes of normalized reward. Both of the two ACC models are calibrated with the microscopic trajectories dataset which is same as the primary parameter training. As we can see from Fig. 10, that the behavior of DRL is relatively decentralized but more aggressive. The average speed of DRL (12.44 m/s) is 3.8% higher than CTH (11.98 m/s), and 7.9% higher than IDM (11.53 m/s). ANOVA analysis shows that the differences of average speed between DRL and either of the ACC models is significant ($P = 0.000$).

Comfort is assessed by jerk in this study. To figure it more clearly, we analyze the distribution of jerks from the extracted 800 episodes. As shown in Fig. 11, the jerk distribution of CTH and IDM are quite similar, while the magnitude of DRL's jerks are much higher. Because the inherent features of analytical model and DDPG are different. A sufficiently small change in the input of continuous function always results in arbitrarily small change in the output. But the relationship between input and output of DRL are instructed by reward function. To explore the influential mechanism of reward setting, we simplify the reward form. Only jerk over threshold are punished, but lower jerk are not rewarded. Although the jerkiness results of DRL varies in a relatively larger range, the general distribution can meet the demand of human ride comfort (no higher than 5.6 m/s^3 , dashed orange line). Only 2.6% of the jerks exceed the critical value.

Therefore, we can conclude that on the premise of relatively comfortable ride, the DRL model is more effective than ACC models. The whole framework of the decision-making training is feasible.

4.6. Extended model integrating CF and LC behavior

Compared to CF, LC behavior is a more challenging driving task, for considering more interactive objects on the road. To testify

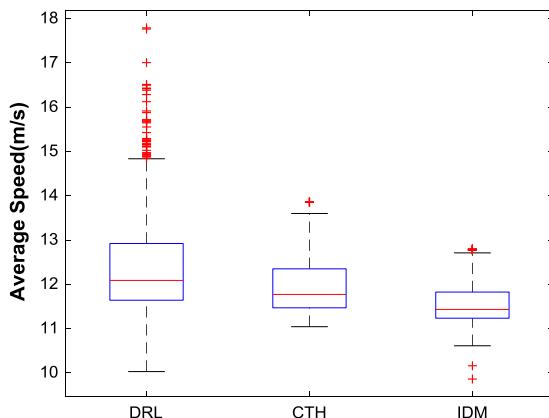


Fig. 10. The average speed distribution of the DRL-CF model and CTH/IDM.

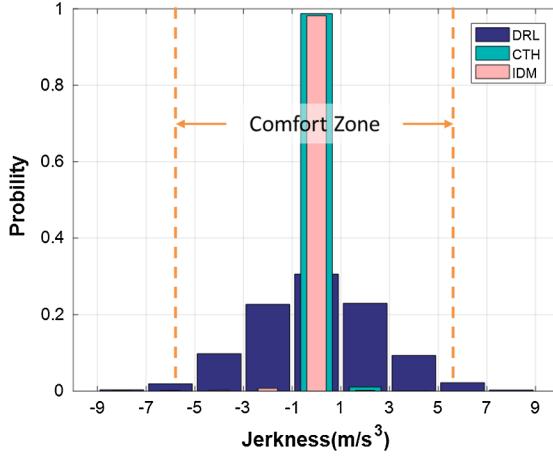


Fig. 11. The jerk distribution of the DRL-CF model and CTH/IDM.

our platform on a more complex environment, and train a model capable of handling most of the traffic scenes, we extended the model integrate CF and LC behavior on a three-lane section. The flow input is set as same as car-following cases, randomly from 600 to 800 veh/h/ln, which is neither heavy nor low for more lane-changing intention. There are two major differences in comparison to CF training:

Input and output of the model. Since LC behavior focuses on not only vehicles on the current lane, but also interactive objects on adjacent lanes. Thus, environment recognition input contains the current speed and acceleration of the AV, the lateral and longitudinal spacing and longitudinal speed information of each nearby vehicle (as shown in Fig. 12, subject vehicle is the blue one, nearby vehicles one downstream and one upstream, 2 lanes on both sides and the current lane are marked). As for movement information output, besides longitudinal acceleration, LC instruction is also provided (3 discrete states: lane-keeping (LK), left lane-changing (LLC), and right lane-changing (RLC)). With a random output of sigmoid activation function, the initial probability of three decisions on LC is identical.

Reward function. LC reward is also incorporated referring to MOBIL model (Kesting et al., 2007), which involves speed of the AV v' , the acceleration increment of the AV and followers on current and subject lane (e.g., if AV in Fig. 12 tends to make left lane-change, the followers are SF and LF). That means, if LC behavior leads sharp brake to followers, even though the AV is accelerated, it will not get a great reward, and even be punished. What's more, when the gap is too small for the AV to cut-in, there will be a punishment (G) and AV will not change the lane. Multiple functions have been tried, the best performed one is listed as below:

$$r = \begin{cases} -1CorR \\ -0.8UorG \\ -0.5L \\ v'car - following \ reward \\ 2 * v' - 0.1 * (\Delta a + \Delta a_{CF} + \Delta a_{LC})lane - changing \ reward \end{cases} \quad (14)$$

In order to simplify the model, on the process of AV's lane-changing, we assume that AV follows the nearer vehicle on both current lane and the target lane for the entire LC process (set to be 2 s in this paper).

4.6.1. Model stability

As shown in Fig. 13, the evaluations on reward, loss and Q value converge over 5000 episodes. It is necessary to explain that the slight fluctuations after convergence are acceptable, which could also be found on other studies (Mnih et al., 2013, 2015).

4.6.2. LC decision analysis

To better understand the improvement of lane-changing decision along the training process, the LC success rate (failed LC

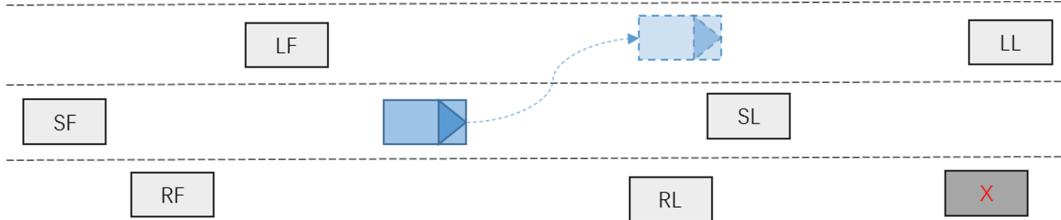


Fig. 12. Description of the nearby vehicles of a LC process.

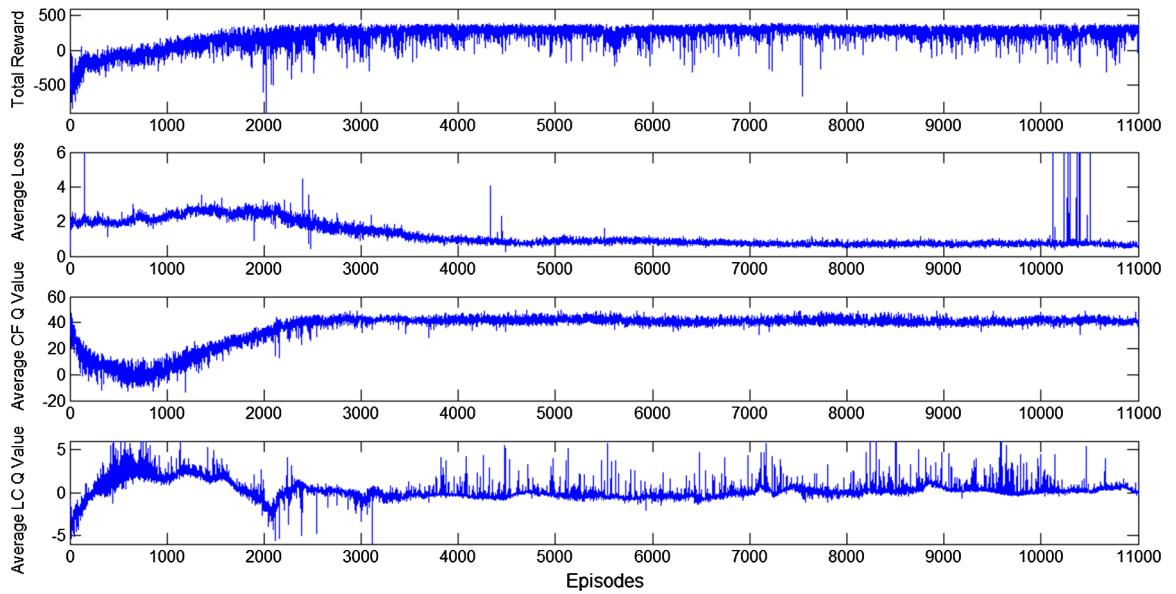


Fig. 13. The evaluation of the integrated model.

happens when the gap is too small even if the LC instruction is given, corresponds to state G) and average reward of successful LC are introduced. We can see from the [Table 3](#), as the number of episodes increases, the frequency of LC instruction within each episode decreases from 267 to 20. The LC success rate, which means the ratio of successful LC decision among all the LC instruction, upgrades firstly and then descends to a stable state around 4%. It indicates that at the beginning, the AV is instructed with LC order frequently (2/3 of the instructions are LC, for the random ratio among LK, LLC and RLC). After thousand episodes of training, it learns not to LC at some wrong moment. What's more, the increase of average reward of success LC, and decrease of actual LC frequency in simulation both further prove the optimization of lane-changing decision. According to our statistics, the LC frequencies of two NGSIM datasets are 0.63 veh/km/ln (US-101) and 0.80 veh/km/ln (I-80), which are similar to our converged result, 0.81 veh/km/ln. However, the traffic flow is heavier in the NGSIM data (around 1600 veh/ln on US-101 and 1300 veh/ln on I-80) comparing our simulation (around 700 veh/ln). Actually, higher LC frequency is desirable since we observe that in some scenarios, the AV is expected to make LC but still follow the front vehicle (see [Fig. 14](#)).

The lead/lag time headways of the LC moment are presented in [Table 4](#). Both the lead time headway and lag headway are significantly higher than the field data of NGSIM dataset. It might result from our LC reward, which considering not only the speed increase of the AV but also the lag vehicles. The objective of our LC model is improving the overall efficiency, while human driver is more willing to get a better driving condition on his own.

4.6.3. Model effectiveness

The average speed of this integrated model (12.74 m/s) is even 2.4% higher than that of the DRL CF model. Therefore, the extended model integrating CF and LC behavior provides a more efficient solution.

In general, the integrated model can provide more efficient driving decision, but the LC behavior so far is still kind of

Table 3
Evaluation of LC training.

Episode	Avg. Freq. of LC instruction	LC success rate (success LC/LC instruction)	Avg. reward of success LC	Actual LC Freq. (veh/km/ln)
< = 50	267	3.2%	2.29	8.62
< = 100	266	3.3%	2.63	8.72
< = 500	235	4.2%	2.86	9.77
500–1000	127	7.5%	3.27	9.43
1000–2000	51	8.0%	3.60	4.06
2000–3000	51	3.5%	4.92	1.82
3000–4000	27	4.4%	5.29	1.18
4000–5000	21	5.6%	5.76	1.19
5000–6000	22	4.7%	5.74	1.02
6000–7000	15	5.9%	6.91	0.89
7000–8000	17	4.5%	6.80	0.77
8000–9000	13	5.8%	6.78	0.73
9000–10000	17	4.0%	5.73	0.66
10000–11000	16	4.3%	6.50	0.69

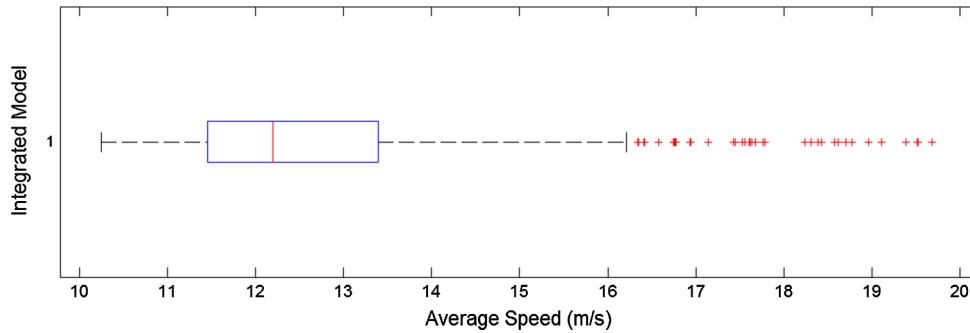


Fig. 14. The average speed distribution of the integrated model.

Table 4
Evaluation of LC gap.

	Avg. lead time headway (s)	Avg. lag time headway (s)
US-101	2.78	3.14
I-80	1.49	1.71
DRL-LC	4.70	4.81

conservative. Therefore, modification on LC reward and thousands of training episodes are required to refine a more rational LC decision, and more specific analysis should be made in our future study.

5. Conclusion

For AVs, subject to the insufficient available data and limitation of the conventional training method, learning an efficient and smart decision-making algorithm seems unachievable. Therefore, we proposed a framework for decision-making learning in this paper, which solve the problems above by introducing the DRL and high-fidelity virtual simulation environment. DDPG and VISSIM are adopted for each of their parts. Firstly, to ensure the convergence, the key element of DRL reward function, is explored for various forms. It concludes that the bounded reward, learning rate and regularized action network should be noted on setting reward. For training acceleration, the primary parameter trained by deep learning and enlarged time step are available. Then, in the aspect of comfort and effectiveness, the trained DRL model is evaluated compared to widely used ACC models IDM and CTH. Results show that the average speed of DRL is 7.9% higher than the IDM, and 3.8% higher than CTH. As for comfort, the ACC models are more comfortable, but only 2.6% of the DRL's jerks exceed the comfort threshold, thus the overall comfort level of DRL model is acceptable. Extending to the more complex task, training a model integrates LC and CF behavior, vehicles perform even better efficiency. The average speed gets 2.4% increase compared to the solely DRL-trained CF model. In consequence, the training framework proposed in this paper for decision making is effective.

However, as shown in the integrated CF and LC case, some minor adjustments are conducted among different scenarios, because the data format and training objective are not always consistent. On the other hand, we consider the training on single AV's decision making in this paper, but the future driving scenario involves cooperative behavior between AVs. Therefore, the investigation of more general model to fit different scenarios and multi-agent decision making will be our future work.

Acknowledgments

This research is sponsored by the National Natural Science Foundation of China (U1764261), the National Key Research and Development Program of China (2018YFB16005), and the “Shuguang Program” supported by Shanghai Education Development Foundation and Shanghai Municipal Education Commission (18SG21). We also would like to thanks PTV to supply the VISSIM Engine for accelerating our DRL training program.

Appendix A. Estimation of the approximated upper bound of average action value

Proposition 1.. *When evaluating the convergence of the current prediction, an upper bound is referred to avoid instability of the commonly used three indexes. The upper bound is expressed as:*

$$\lim_{t \rightarrow \infty} Q_t \approx \frac{r}{1 - \gamma} \quad (\text{A1})$$

Proof.. According to Eq.(1), we expand the action value as:

$$\begin{aligned}
Q^\pi(s, a) &= \mathbb{E}(r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s, a) \\
&= \mathbb{E}(r_s^a + \gamma Q^\pi(s', a') | s, a) \\
&\approx r_s^a + \gamma Q^\pi(s', a')
\end{aligned} \tag{A2}$$

Denote

$$Q_t = Q^\pi(s, a) \tag{A3}$$

$$Q_{t+1} = Q^\pi(s', a') \tag{A4}$$

$$r_t = r_s^a \tag{A5}$$

$$r = \max(r_t) \tag{A6}$$

So, we have

$$Q_t \approx r_t + \gamma Q_{t+1} \tag{A7}$$

When the training converges, every step is deemed to have maximized reward. Thus for any step t ,

$$Q_t \approx r + \gamma Q_{t+1} \tag{A8}$$

$$Q_t \approx Q_{t+1} \tag{A9}$$

Then, we have the approximated upper bound of average action value:

$$\lim_{t \rightarrow \infty} Q_t \approx \frac{r}{1 - \gamma} \tag{A10}$$

References

- Bojarski, M., Del Testa, D., Dworakowski, D., Firner, B., Flepp, B., Goyal, P., Zhang, X., 2016. End to end learning for self-driving cars. arXiv preprint arXiv:1604.07316.
- Buckland, M., Interview with Jeff Hannan," Publication date unknown. Available: <<http://www.aijunkie.com/misc/hannan/hannan.html>> .
- Buehler, M., Iagnemma, K., Singh, S. (Eds.). (2009). The DARPA urban challenge: autonomous vehicles in city traffic, vol. 56. Springer.
- Cardamone, L., Loiacono, D., Lanzi, P.L., 2009. On-line neuroevolution applied to the open racing car simulator. In: IEEE Congress on Evolutionary Computation, 2009. CEC'09. IEEE, pp. 2622–2629.
- Cardamone, L., Loiacono, D., Lanzi, P.L., 2009. Evolving competitive car controllers for racing games with neuroevolution. In: Proceedings of the 11th Annual Conference on Genetic and Evolutionary Computation, ACM, pp. 1179–1186.
- Chen, C., Seff, A., Kornhauser, A., Xiao, J., 2015. Deepdriving: Learning affordance for direct perception in autonomous driving. In: Proceedings of the IEEE International Conference on Computer Vision, pp. 2722–2730.
- Chen, X.W., Lin, X., 2014. Big data deep learning: challenges and perspectives. *IEEE Access* 2, 514–525.
- De La Escalera, A., Moreno, L.E., Salichs, M.A., Armingol, J.M., 1997. Road traffic sign detection and classification. *IEEE Trans. Ind. Electron.* 44 (6), 848–859.
- Deng, J., Dong, W., Socher, R., Li, L.J., Li, K., Fei-Fei, L., 2009. Imagenet: A large-scale hierarchical image database. In: IEEE Conference on Computer Vision and Pattern Recognition, 2009. CVPR 2009. IEEE, pp. 248–255.
- Farabet, C., Couprie, C., Najman, L., LeCun, Y., 2013. Learning hierarchical features for scene labeling. *IEEE Trans. Pattern Anal. Mach. Intell.* 35 (8), 1915–1929.
- Fritsch, J., Kuhl, T., Geiger, A., 2013. A new performance measure and evaluation benchmark for road detection algorithms. In: 2013 16th International IEEE Conference on Intelligent Transportation Systems-(ITSC). IEEE, pp. 1693–1700.
- Geiger, A., Lenz, P., Stiller, C., Urtasun, R., 2013. Vision meets robotics: The KITTI dataset. *Int. J. Robot. Res.* 32 (11), 1231–1237.
- Henderson, P., Islam, R., Bachman, P., Pineau, J., Precup, D., Meger, D., 2017. Deep Reinforcement Learning that Matters.
- Hinton, G.E., Osindero, S., Teh, Y.W., 2006. A fast learning algorithm for deep belief nets. *Neural Comput.* 18 (7), 1527–1554.
- Ilyas, A., Engstrom, L., Santurkar, S., Tsipras, D., Janoos, F., Rudolph, L., Madry, A., 2018. Are Deep Policy Gradient Algorithms Truly Policy Gradient Algorithms?
- Islam, R., Henderson, P., Gomrokchi, M., Precup, D., 2017. Reproducibility of Benchmarked Deep Reinforcement Learning Tasks for Continuous Control.
- Kesting, A., Treiber, M., Helbing, D., 2007a. Mobil: general lane-changing model for car-following model. *Transp. Res. Rec. J. Transp. Res. Board* 1999 (1), 86–94.
- Kesting, A., Treiber, M., Schönhof, M., Helbing, D., 2007b. Extending adaptive cruise control to adaptive driving strategies. *Transp. Res. Rec. J. Transp. Res. Board* 2000, 16–24.
- Kesting, A., Treiber, M., Schönhof, M., Helbing, D., 2008. Adaptive cruise control design for active congestion avoidance. *Transp. Res. Part C: Emerg. Technol.* 16 (6), 668–683.
- Kim, K.J., Seo, J.H., Park, J.G., Na, J.C., 2012. Generalization of TORCS car racing controllers with artificial neural networks and linear regression analysis. *Neurocomputing* 88, 87–99.
- Kober, J., Peters, J., 2010. Imitation and reinforcement learning. *IEEE Rob. Autom. Mag.* 17 (2), 55–62.
- Koutnik, J., Schmidhuber, J., Gomez, F., 2014. Evolving deep unsupervised convolutional networks for vision-based reinforcement learning. In: Proceedings of the 2014 Annual Conference on Genetic and Evolutionary Computation. ACM, pp. 541–548.
- Lau, Y. P., 2016. Using keras and deep deterministic policy gradient to play torcs. Available: <<https://yanpanlau.github.io/2016/10/11/Torcs-Keras.html>> .
- LeCun, Y., Bottou, L., Bengio, Y., Haffner, P., 1998. Gradient-based learning applied to document recognition. *Proc. IEEE* 86 (11), 2278–2324.
- Li, P.Y., Shrivastava, A., 2002. Traffic flow stability induced by constant time headway policy for adaptive cruise control vehicles. *Transp. Res. Part C: Emerg. Technol.* 10 (4), 275–301.
- Lillicrap, T.P., Hunt, J.J., Pritzel, A., et al., 2015. Continuous control with deep reinforcement learning. arXiv preprint arXiv:1509.02971.
- Loiacono, D., Prete, A., Lanzi, P.L., Cardamone, L., 2010. Learning to overtake in TORCS using simple reinforcement learning. In: Evolutionary Computation. IEEE, pp. 1–8.
- Maas, A.L., Hannun, A.Y., Ng, A.Y., 2013. Rectifier nonlinearities improve neural network acoustic models. In: ICML '13. p. 6.
- Mania, H., Guy, A., Recht, B., 2018. Simple random search provides a competitive approach to reinforcement learning.
- McCall, J.C., Trivedi, M.M., 2004. An integrated, robust approach to lane marking detection and lane tracking. In: Intelligent Vehicles Symposium, 2004 IEEE. IEEE,

- pp. 533–537.
- Mnih, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D., et al., 2013. Playing Atari with deep reinforcement learning. arXiv preprint arXiv:1312.5602 .
- Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A.A., Veness, J., Bellemare, M.G., et al., 2015. Human-level control through deep reinforcement learning. *Nature* 518 (7540), 529–533.
- Mnih, V., Badia, A.P., Mirza, M., Graves, A., Lillicrap, T., Harley, T., Kavukcuoglu, K., 2016. June). Asynchronous methods for deep reinforcement learning. In: International Conference on Machine Learning, pp. 1928–1937.
- Najafabadi, M.M., Villanustre, F., Khoshoftaar, T.M., Seliya, N., Wald, R., Muharemagic, E., 2015. Deep learning applications and challenges in big data analytics. *J. Big Data* 2 (1), 1.
- Perot, E., Jaritz, M., Toromanoff, M., Charette, R.D., 2017. End-to-End Driving in a Realistic Racing Game with Deep Reinforcement Learning. In: IEEE Conference on Computer Vision and Pattern Recognition Workshops. IEEE Computer Society, pp. 474–475.
- Richter, S.R., Vineet, V., Roth, S., Koltun, V., 2016. Playing for data: Ground truth from computer games. In: European Conference on Computer Vision. Springer International Publishing, pp. 102–118.
- Sallab, A.E., Abdou, M., Perot, E., Yogamani, S., 2017. Deep reinforcement learning framework for autonomous driving. *Electron. Imag.* 2017 (19), 70–76.
- Sermanet, P., LeCun, Y., 2011. Traffic sign recognition with multi-scale convolutional networks. In: The 2011 International Joint Conference on Neural Networks (IJCNN). IEEE, pp. 2809–2813.
- Swaroop, D., Rajagopal, K.R., 2001. A review of constant time headway policy for automatic vehicle following. In: Intelligent Transportation Systems, 2001. Proceedings. 2001 IEEE. IEEE, pp. 65–69.
- Szarvas, M., Yoshizawa, A., Yamamoto, M., Ogata, J., 2005. Pedestrian detection with convolutional neural networks. In: Intelligent vehicles symposium, 2005. Proceedings. IEEE. IEEE, pp. 224–229.
- Teaching, E.S., 2018. CS231n Convolutional Neural Networks for Visual Recognition [WWW Document]. <<http://vision.stanford.edu/teaching/cs231n/>> .
- Togelius, J., Nardi, R.D., Lucas, S.M., 2006. Making racing fun through player modeling and track evolution. In: Proceedings of the Sab'06 Workshop on Adaptive Approaches for Optimizing Player Satisfaction in Computer & Physical Games, pp. 70.
- Togelius, J., Nardi, R.D., Lucas, S.M., 2007. Towards automatic personalized content creation for racing games. In: IEEE Symposium on Computational Intelligence and Games. IEEE Computer Society, pp.252–259.
- Torcs, the open racing car simulator. 2017. Available: <<http://torcs.sourceforge.net/>> .
- Van Hoorn, N., Togelius, J., Wierstra, D., Schmidhuber, J., 2009. Robust player imitation using multiobjective evolution. In: IEEE Congress on Evolutionary Computation, 2009. CEC'09. IEEE, pp. 652–659.
- Van Winsum, W., 1999. The human element in car following models. *Transp. Res. Part F: Traffic Psychol. Behav.* 2 (4), 207–211.
- Wang, H., Cai, Y., Chen, L., 2014. A vehicle detection algorithm based on deep belief network. *The scientific world journal*, 2014.
- Watkins, C.J.C.H., 1989. Learning from delayed rewards (Doctoral dissertation, University of Cambridge).
- Wei, X., 2015. Must know tips/tricks in deep neural network. Available: <<http://lamda.nju.edu.cn/weixs/project/CNNTricks/CNNTricks.html>> .